

COMS4047A - Reinforcement Learning

Lab 3 - Deep Q-Network

Shahil Mawjee

Overview

The lab will focus on the implementation of the Deep Q-Learning (DQN) agent to play Atari games. For this lab we will use the Pong ([link](#)) environment which are contained within the OpenAI Gym **Atari** package.

Goals:

- Implement model from research paper.
- Understand hyperparameter optimisation.

OpenAI Gym Wrapper

Wrappers are used to transform an environment in a modular way. [Link to gym wrappers](#).

```
env = gym.make('PongNoFrameskip-v4')  
env = MyWrapper(env)
```

WarpFrame Wrapper

Warp frames to 84x84 as done in the Nature paper and later work. Expects inputs to be of shape height x width x number_of_channels

PyTorchFrame Wrapper

Pytorch expects images in the form [number_of_channels, height, width] hence this wrapper transforms image from [height, width, number_of_channels] to [number_of_channels, height, width]

1 Deep Q-Learning (DQN)

Refer to the following papers in order to implement the DQN:

- Human-level control through deep reinforcement learning (link)
- Playing Atari with Deep Reinforcement Learning (link)




Figure 1: Deep Q Network (Source: Nature paper)

Algorithm 1 Deep Q-learning with Experience Replay

```

Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$ 
  end for
end for

```

1.1 Environment: Atari Pong

For this lab we will use Atari's Pong game ([link](#)). In this environment, the observation is an RGB image of the screen, which is an array of shape (210, 160, 3)




Figure 2: Pong Game (Source - OpenAI Gym website)

1.2 Training Loop

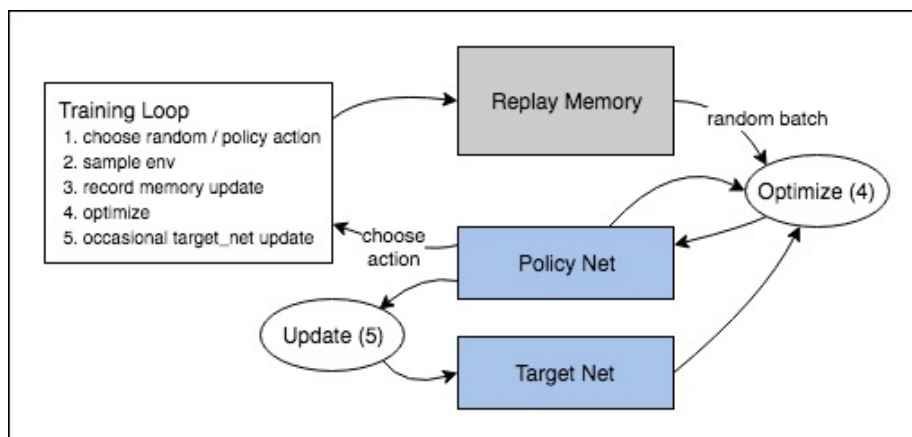


Figure 3: Training loop overview

Submission

Include the following in your zip file:

- Summary of process write up (half a page)
- Your code
- Plot of reward per episode (as done in the paper)
- GIF of your trained agent playing Pong

Submit your zip file to Moodle.